



OpenStack ICE-house Deployment on Ubuntu Server 14.04 LTS Server

Draft - 1.0

This configuration guide includes Multi-node OpenStack deployment on Ubuntu Server 14.04 LTS (server) platform. In my case Ubuntu servers are nested in VMware ESXi 5.5.

Setup requirements are as follows,

- 1 ESXi Server with minimum 64 GB memory, 1 TB hard drive and Intel cpu with min 8 cores
- 3 Ubuntu 14.04 LTS server edition VMs with Multiple NICs. Controller and Compute node with 2 NICs each and Network node with 3 NICs
- Internet connectivity



OpenStack deployment is also possible on ESXi server having lower memory. Example ESXi server with 32 Gb memory and 4 cpus however it is always better to have esxi host with 64GB Memory

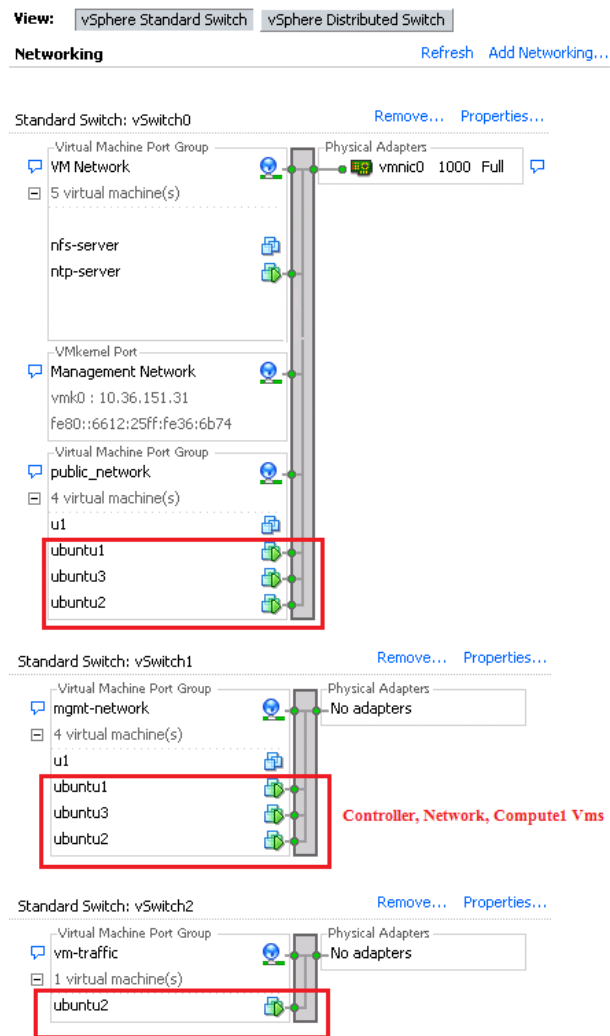
Let's begin,

Download Ubuntu 14.04 LTS Server ISO from Internet and copy it on the data store of the ESX host on which Openstack host vms will get deployed.

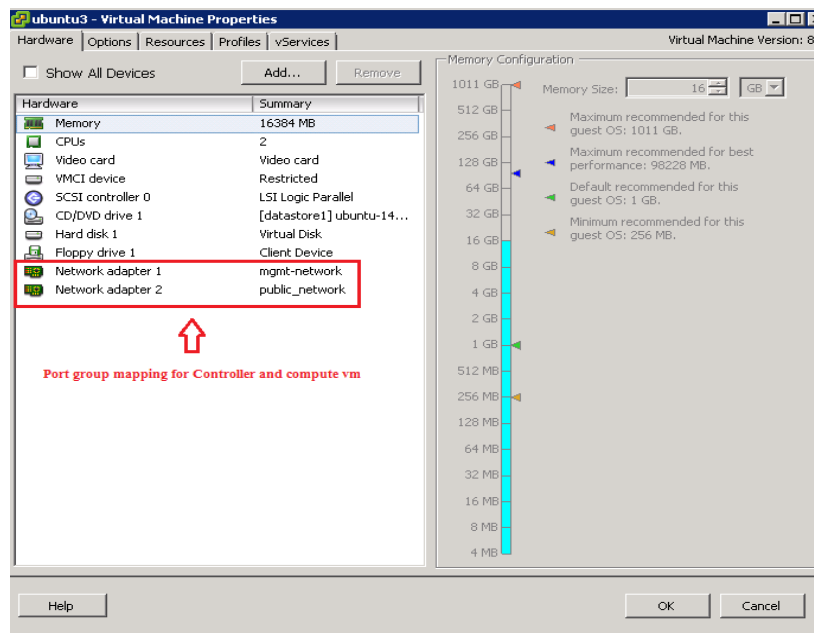
You may choose the latest Ubuntu version if any. Currently 14.0.4 LTS server is the latest version hence I will recommend to use the same version.

- **Create Controller VM** with 16GB memory, 100GB Hard drive, 2 NICs with VMXnet3 Adaptor type, min 2 vCpus with one core each and Install Ubuntu.
- **Create Network VM** with 12GB memory, 100GB Hard drive, 3 NICs with VMXnet3 Adaptor type, min 2 vCpus with one core each and install Ubuntu
- **Create Compute VM** with 16GB memory, 100GB Hard drive, 2 NICs with VMXnet3 Adaptor type, min 2 vCpus with one core each and install Ubuntu

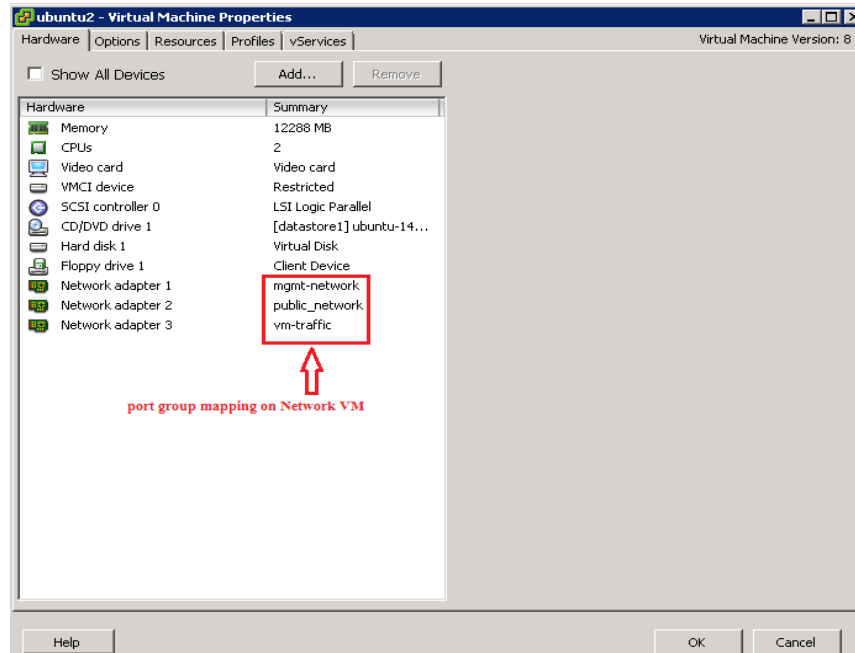
Refer ESXi server network configuration port group mapping



Port group mapping screenshots for controller and compute



Port group mapping screenshots for controller and compute



Ubuntu Installation Instructions

- Once the above blank vms are ready with the port group mapping, begin Ubuntu installation
- Select open-ssh server package on package selection menu
- Proceed with default installation options

Network Topology -

I have assigned the following IPs. User can use any IP pool however corporate / Public network should be available for internet access

Management Network - All openstack components will communicate using this network. In my case management network is 10.0.0.x/24

Corp/ Public Network - This network has internet connectivity. The IPs from this network will be used for the following operations

Refer the following IP configuration chart

Node Type	Interface	Network Connection	IP address
Controller	Eth0	Management	10.0.0.11 /24
	Eth1	Corporate/ Public	10.36.30.83/16
Network	Eth0	Management	10.0.0.21/24
	Eth1	Corporate/ Public	10.36.30.84/16
	Eth2	VM-traffic	20.0.0.31
Compute1	Eth0	Management	10.0.0.31/24
	Eth1	Corporate/ Public	10.36.30.84/16
	Eth2	VM-traffic	20.0.0.41



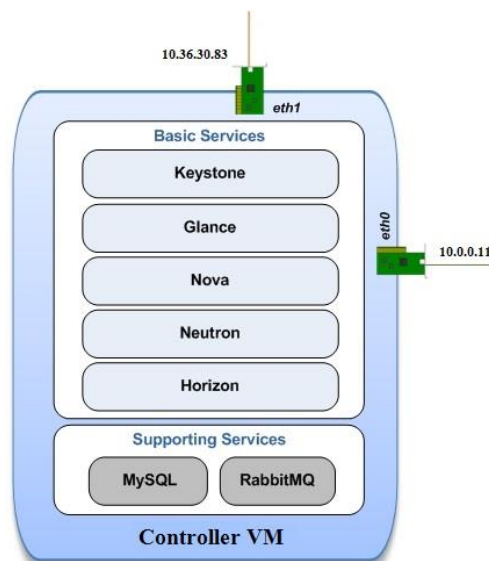
Note

Corporate/ Public Network has internet connectivity and it will be used to ssh Ubuntu VMs. Internet connection is also necessary to download OpenStack packages from Ubuntu repository.

Compute node is named as compute1 because it is possible to add more compute nodes as per the requirement. It is also possible to add multiple Network nodes.

Configure Controller Node

Here is a small diagram of basic services (keystone, glance, nova, neutron and horizon) that are installed on the Controller VM and also the supporting services such as MySQL database, message broker (RabbitMQ), and NTP.



Let's Start preparing Controller VM

Following tasks will be performed in the following section,

- Upgrade Ubuntu OS along with repository update to get the latest package repos
- Install the NTP daemon
- Install MySQL server and configure
- Install RabbitMQ server
- Install and configure Keystone

▪ Update and Upgrade your System

```
apt-get update -y && apt-get upgrade -y && apt-get dist-upgrade
```

▪ Install NTP service (Network Time Protocol)

```
apt-get install -y ntp
```

▪ Install MySQL

```
apt-get install -y mysql-server python-mysqldb
```

- Set the bind-address key to the management IP address of the controller node

```
vi /etc/mysql/my.cnf
bind-address = 10.0.0.11
```

- Under the [mysqld] section, set the following keys to enable InnoDB, UTF-8 character set, and UTF-8 collation by default

```
vi /etc/mysql/my.cnf
[mysqld]
default-storage-engine = innodb
innodb_file_per_table
collation-server = utf8_general_ci
init-connect = 'SET NAMES utf8'
character-set-server = utf8
```

- Restart the MySQL service:

```
service mysql restart
```

- Delete the anonymous users that are created when the database is first started

```
mysql_install_db
mysql_secure_installation
```



Component Introduction - RabbitMQ

- OpenStack requires a messaging service for internal communications. Between various components Messaging enables software applications to connect and scale. Applications can connect to each other, as components of a larger application, or to user devices and data. Typically RabbitMQ is used in Openstack environment.
- Default RabbitMQ port is 5672

- Install RabbitMQ (Message Queue)

```
apt-get install -y rabbitmq-server
```

Keystone installation and configuration



Component Introduction - **Keystone** (Identity service)

- Keystone is a framework for authentication and authorization for all the OpenStack services.
- Keystone handles API requests as well as providing configurable catalog, policy, token and identity services.
- It provides the ability to add users to groups (also known as tenants) and to manage permissions between users and groups. Permissions include the ability to launch and terminate instances.

■ Install keystone packages

```
apt-get install -y keystone
```

■ Create a MySQL database for keystone

```
mysql -u root -p

CREATE DATABASE keystone;
GRANT ALL PRIVILEGES ON keystone.* TO 'keystone'@'localhost' IDENTIFIED BY
'KEYSTONE_DBPASS';
GRANT ALL PRIVILEGES ON keystone.* TO 'keystone'@'%' IDENTIFIED BY
'KEYSTONE_DBPASS';

exit;
```



'localhost' can be replaced by controller management ip 10.0.0.11

■ Remove Keystone SQLite database

```
rm /var/lib/keystone/keystone.db
```

■ Edit /etc/keystone/keystone.conf

```
vi /etc/keystone/keystone.conf

[database]
replace connection = sqlite:///var/lib/keystone/keystone.db by
connection = mysql://keystone:KEYSTONE_DBPASS@10.0.0.11/keystone

[DEFAULT]
```



```
admin token=ADMIN
log_dir=/var/log/keystone
```

- Restart the identity service then synchronize the database

```
service keystone restart
keystone-manage db_sync
```

- Check synchronization

```
mysql -u root -p keystone
show TABLES;
```

- Define users, tenants, and roles

```
export OS_SERVICE_TOKEN=ADMIN
export OS_SERVICE_ENDPOINT=http://10.0.0.11:35357/v2.0
```

```
#Create an administrative user
keystone user-create --name=admin --pass=admin_pass --email=admin@domain.com
keystone role-create --name=admin
keystone tenant-create --name=admin --description="Admin Tenant"
keystone user-role-add --user=admin --tenant=admin --role=admin
keystone user-role-add --user=admin --role=_member_ --tenant=admin

#Create a normal user
keystone user-create --name=demo --pass=demo_pass --email=demo@domain.com
keystone tenant-create --name=demo --description="Demo Tenant"
keystone user-role-add --user=demo --role=_member_ --tenant=demo

#Create a service tenant
keystone tenant-create --name=service --description="Service Tenant"
```

- Define services and API endpoints

```
keystone service-create --name=keystone --type=identity --
description="OpenStack Identity"

keystone endpoint-create --service-id(service id which was generated by the above
command) --publicurl=http://10.36.30.83:5000/v2.0
internalurl=http://10.0.0.11:5000/v2.0
--adminurl=http://10.0.0.11:35357/v2.0
```



for example, keystone endpoint-create --service-id15c11a23667e427e91bc31335b45f4bd publicurl=http://10.36.30.83:5000/v2.0 internalurl=http://10.0.0.11:5000/v2.0 --adminurl=http://10.0.0.11:35357/v2.0

- Create a simple credential file

```
vi creds
#Paste the following:
export OS_TENANT_NAME=admin
export OS_USERNAME=admin
export OS_PASSWORD=admin_pass
```

```
export OS_AUTH_URL=http://10.36.30.83:5000/v2.0/
```

```
vi admin_creds
#Paste the following:
export OS_USERNAME=admin
export OS_PASSWORD=admin_pass
export OS_TENANT_NAME=admin
export OS_AUTH_URL=http://10.0.0.11:35357/v2.0
```

Test Keystone

- Clear the values in the `OS_SERVICE_TOKEN` and `OS_SERVICE_ENDPOINT` environment variables

```
unset OS_SERVICE_TOKEN OS_SERVICE_ENDPOINT
```

- Request a authentication token

```
keystone --os-username=admin --os-password=admin pass --os-auth-
url=http://10.0.0.11:35357/v2.0 token-get
```

- Load credential admin file

```
source admin_creds
keystone token-get
```

- Load credential file

```
source creds
keystone user-list
keystone user-role-list --user admin --tenant admin
```

Glance installation and configuration



Component Introduction - [Glance](#)

- The OpenStack Image Service called Glance provides discovery, registration and delivery services for disk and server images.
- It has the ability to copy (or snapshot) a server image and then to store it promptly. Stored images then can be used as templates to get new servers up and running quickly, and can also be used to store and catalog unlimited backups.
- Virtual-machine images can be stored in various locations, including simple file systems and object-storage systems such as OpenStack Object Storage (code named Swift).
- Glance includes **Glance-api**, **Glance-registry**, **Database**, **Storage repository**

- **Install Glance packages**

```
apt-get install -y glance python-glanceclient
```

- **Create a MySQL database for Glance**

```
mysql -u root -p

CREATE DATABASE glance;

GRANT ALL PRIVILEGES ON glance.* TO 'glance'@'localhost' IDENTIFIED BY
GLANCE_DBPASS';

GRANT ALL PRIVILEGES ON glance.* TO 'glance'@'%' IDENTIFIED BY
'GLANCE_DBPASS';

exit;
```

- **Configure service user and role**

```
keystone user-create --name=glance --pass=service_pass --email=glance@domain.com

keystone user-role-add --user=glance --tenant=service --role=admin
```

- **Register the service and create the endpoint**

```
keystone service-create --name=glance --type=image --description="OpenStack Image Service"

keystone endpoint-create --service-id= (service id which was generated by the above command)

--publicurl=http://10.36.30.83:9292 --internalurl=http://10.0.0.11:9292 --adminurl=http://10.0.0.11:9292
```



*for example, keystone endpoint-create --service-id15c11a23667e427e91bc31335b45f4bd
--publicurl=http://10.36.30.83:9292 --internalurl=http://10.0.0.11:9292 --adminurl=http://10.0.0.11:9292*

- **Update /etc/glance/glance-api.conf**

```
vi /etc/glance/glance-api.conf

[database]

replace sqlite_db = /var/lib/glance/glance.sqlite with
```

```

connection = mysql://glance:GLANCE_DBPASS@10.0.0.11/glance

[DEFAULT]

rpc_backend = rabbit
rabbit_host = 10.0.0.11

[keystone_authtoken]
auth_uri = http://10.0.0.11:5000
auth_host = 10.0.0.11
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = glance
admin_password = service_pass

[paste_deploy]
flavor = keystone

```

- **Update /etc/glance/glance-registry.conf**

```

vi /etc/glance/glance-registry.conf

[database]
replace sqlite_db = /var/lib/glance/glance.sqlite with:
connection = mysql://glance:GLANCE_DBPASS@10.0.0.11/glance

[keystone_authtoken]
auth_uri = http://10.0.0.11:5000
auth_host = 10.0.0.11
auth_port = 35357
auth_protocol = http
admin_tenant_name = service

```

```
admin_user = glance
admin_password = service_pass

[paste_deploy]

flavor = keystone
```

- Restart the glance-api and glance-registry services

```
service glance-api restart; service glance-registry restart
```

- Synchronize the glance database

```
glance-manage db_sync
```

- Test Glance, upload the cirros cloud image

```
source creds

glance image-create --name "cirros-0.3.2-x86_64" --is-public true \
--container-format bare --disk-format qcow2 \
--location http://cdn.download.cirros-cloud.net/0.3.2/cirros-0.3.2-x86_64-
disk.img
```

```
glance image-list
```

Install the compute Service (Nova)



Component Introduction – Nova

- Nova, also known as OpenStack Compute, is the software that controls your Infrastructure as a Service (IaaS) cloud computing platform. It is similar in scope to Amazon EC2 and Rackspace Cloud Servers. Nova does not include any virtualization software, rather it defines drivers that interact with underlying virtualization mechanisms that run on your host operating system, and exposes functionality over a web API.
- Nova has a concept of Fixed IPs and Floating IPs. Fixed IPs are assigned to an instance on creation and stay the same until the instance is explicitly terminated. Floating IPs are IP addresses that can be dynamically associated with an instance. This address can be disassociated and associated with another instance at any time.

- Install nova packages

```
apt-get install -y nova-api nova-cert nova-conductor nova-consoleauth \
nova-novncproxy nova-scheduler python-novaclient
```

- Create a Mysql database for Nova

```
mysql -u root -p

CREATE DATABASE nova;

GRANT ALL PRIVILEGES ON nova.* TO 'nova'@'localhost' IDENTIFIED BY
'NOVA_DBPASS';

GRANT ALL PRIVILEGES ON nova.* TO 'nova'@'%' IDENTIFIED BY 'NOVA_DBPASS';

exit;
```

- Configure service user and role

```
keystone user-create --name=nova --pass=service_pass --email=nova@domain.com

keystone user-role-add --user=nova --tenant=service --role=admin
```

- Register the service and create the endpoint

```
keystone service-create --name=nova --type=compute --description="OpenStack
Compute"

keystone endpoint-create --service-id= (service id which was generated by the above command) --
publicurl=http://10.36.30.83:8774/v2/%(tenant_id)s --
internalurl=http://10.0.0.11:8774/v2/%(tenant_id)s --
adminurl=http://10.0.0.11:8774/v2/%(tenant_id)s
```



*for example, keystone endpoint-create --service-id15c11a23667e427e91bc31335b45f4bd
--publicurl=http://10.36.30.83:8774/v2/%(tenant_id)s --internalurl=http://10.0.0.11:8774/v2/%(tenant_id)s
--adminurl=http://10.0.0.11:8774/v2/%(tenant_id)s*

- Edit the `/etc/nova/nova.conf`

```
vi /etc/nova/nova.conf

[database]
connection = mysql://nova:NOVA_DBPASS@controller/nova

[DEFAULT]
rpc_backend = rabbit
rabbit_host = controller
my_ip = 10.0.0.11
vncserver_listen = 10.0.0.11
vncserver_proxyclient_address = 10.0.0.11
auth_strategy = keystone

[keystone_authtoken]
auth_uri = http://10.0.0.11:5000
auth_host = 10.0.0.11
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = nova
admin_password = service_pass
```

- Remove Nova SQLite database

```
rm /var/lib/nova/nova.sqlite
```

- **Synchronize your database**

```
nova-manage db sync
```

- **Restart nova-* services**

```
service nova-api restart  
service nova-cert restart  
service nova-conductor restart  
service nova-consoleauth restart  
service nova-novncproxy restart  
service nova-scheduler restart
```

- **nova-manage service list**

To verify your configuration, list available images

```
source creds  
nova image-list
```

Install Neutron



Component Introduction – Neutron

- Neutron(formerly Quantum) is a system for managing networks and IP addresses. Like other aspects of the cloud operating system, it can be used by administrators and users to increase the value of existing data center assets. Neutron ensures the network will not be the bottleneck or limiting factor in a cloud deployment and gives users real self-service, even over their network configurations.
- Neutron provides networking models for different applications or user groups. Standard models include flat networks or VLANs for separation of servers and traffic. Neutron manages IP addresses, allowing for dedicated static IPs or DHCP. Floating IPs allow traffic to be dynamically re routed to any of your compute resources, which allows you to redirect traffic during maintenance or in the case of failure.
- Users can create their own networks, control traffic and connect servers and devices to one or more networks.

- Install the Neutron server and the OpenVSwitch packages

```
apt-get install -y neutron-server neutron-plugin-ml2
```

- Create a MySql database for Neutron

```
mysql -u root -p

CREATE DATABASE neutron;

GRANT ALL PRIVILEGES ON neutron.* TO neutron@'localhost' IDENTIFIED BY
'NEUTRON_DBPASS';

GRANT ALL PRIVILEGES ON neutron.* TO neutron@'%' IDENTIFIED BY
'NEUTRON_DBPASS';

exit;
```

- Configure service user and role

```
keystone user-create --name=neutron --pass=service_pass --
email=neutron@domain.com

keystone user-role-add --user=neutron --tenant=service --role=admin
```

- Register the service and create the endpoint

```
keystone service-create --name=neutron --type=network --
description="OpenStack Networking"

keystone endpoint-create --service-id=(service id which was generated by the above command)

--publicurl=http://10.36.30.83:9696 --internalurl=http://10.0.0.11:9696

--adminurl=http://10.0.0.11:9696
```



*for example, keystone endpoint-create --service-id15c11a23667e427e91bc31335b45f4bd
--publicurl=http://10.36.30.83:9696 --internalurl=http://10.0.0.11:9696 --
adminurl=http://10.0.0.11:9696*

- Update /etc/neutron/neutron.conf

```
vi /etc/neutron/neutron.conf

[database]
```

```

replace connection = sqlite:///var/lib/neutron/neutron.sqlite with
connection = mysql://neutron:NEUTRON_DBPASS@10.0.0.11/neutron

[DEFAULT]

replace core_plugin = neutron.plugins.ml2.plugin.Ml2Plugin with
core_plugin = ml2

service_plugins = router

allow_overlapping_ips = True

auth_strategy = keystone

rpc_backend = neutron.openstack.common.rpc.impl_kombu

rabbit_host = 10.0.0.11

notify_nova_on_port_status_changes = True

notify_nova_on_port_data_changes = True

nova_url = http://10.0.0.11:8774/v2

nova_admin_username = nova

# with the output of this command (keystone tenant-list | awk '/ service / {
print $2 }')

nova_admin_tenant_id = (put SERVICE_TENANT_ID here)

#service tenant id can be found with the help of this command "keystone
tenant-list"

nova_admin_password = service_pass

nova_admin_auth_url = http://10.0.0.11:35357/v2.0


[keystone_authtoken]

auth_uri = http://10.0.0.11:5000

auth_host = 10.0.0.11

auth_port = 35357

auth_protocol = http

admin_tenant_name = service

admin_user = neutron

admin_password = service_pass

```



If you get an error message "Expecting an auth URL via either --os-auth-url or env[OS_AUTH_URL]" after typing keystone tenant-list command, you probably need to make sure that the /cred file is updated correctly.

Following are the cred file entries

```
export OS_TENANT_NAME=admin
export OS_USERNAME=admin
export OS_PASSWORD=admin_pass
export OS_AUTH_URL="http://192.168.100.51:5000/v2.0/"d type "source creds"
```

Configure Modular Layer 2 (ML2) plug-in



Component Introduction – ML2 Plugin for Neutron

- Starting with Havana release, openvswitch and linuxbridge plugins are deprecated. Modular Layer 2 (ML2) plugin replaces these plugins. L2 agents work with ML2 plugin and continue to work with the deprecated monolithic plugins.

■ Edit ml2_conf.ini

```
vi /etc/neutron/plugins/ml2/ml2_conf.ini

[ml2]

type_drivers = gre

tenant_network_types = gre

mechanism_drivers = openvswitch


[ml2_type_gre]

tunnel_id_ranges = 1:1000


[securitygroup]

firewall_driver = neutron.agent.linux.iptables_firewall.OVSHybridIptablesFirewallDriver

enable_security_group = True
```

- **Configure Compute to use Networking**

add in `/etc/nova/nova.conf`

```
vi /etc/nova/nova.conf

[DEFAULT]

network_api_class=nova.network.neutronv2.api.API

neutron_url=http://10.0.0.11:9696

neutron_auth_strategy=keystone

neutron_admin_tenant_name=service

neutron_admin_username=neutron

neutron_admin_password=service_pass

neutron_admin_auth_url=http://10.0.0.11:35357/v2.0

libvirt_vif_driver=nova.virt.libvirt.vif.LibvirtHybridOVSBridgeDriver

linuxnet_interface_driver=nova.network.linux_net.LinuxOVSInterfaceDriver

firewall_driver=nova.virt.firewall.NoopFirewallDriver

security_group_api=neutron
```

- **Restart the Compute services**

```
service nova-api restart

service nova-scheduler restart

service nova-conductor restart
```

- **Restart the Networking service**

```
service neutron-server restart
```

Install the dashboard Service (Horizon)



Component Introduction - Horizon

- Horizon is a Django-based project aimed at providing a complete OpenStack Dashboard.
- Horizon is the canonical implementation of OpenStack's Dashboard, which provides a web based user interface to OpenStack services including Nova, Swift etc.

■ Install the required packages

```
apt-get install -y apache2 memcached libapache2-mod-wsgi openstack-dashboard
```

■ You can remove the openstack-dashboard-ubuntu-theme package

```
apt-get remove -y --purge openstack-dashboard-ubuntu-theme
```

■ Edit /etc/openstack-dashboard/local_settings.py

```
vi /etc/openstack-dashboard/local_settings.py
ALLOWED_HOSTS = ['localhost', '10.36.30.83']
OPENSTACK_HOST = "10.0.0.11"
```

■ Reload Apache and memcached

```
service apache2 restart; service memcached restart
```



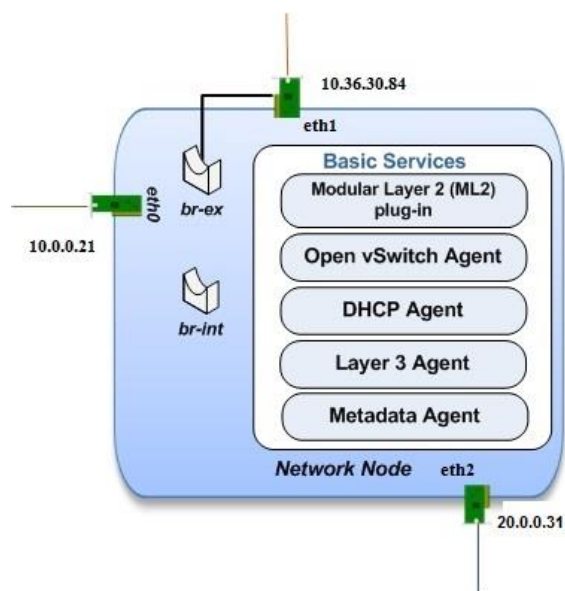
If you have this error: apache2: Could not reliably determine the server's fully qualified domain name, using 127.0.1.1. Set the 'ServerName' directive globally to suppress this message"

Solution: Edit /etc/apache2/apache2.conf

```
vi /etc/apache2/apache2.conf
Add the following new line end of file:
ServerName localhost
```

Reload Apache and memcached

Configure Network Node



Update and Upgrade your System

```
apt-get update -y && apt-get upgrade -y && apt-get dist-upgrade
```

Install NTP service

```
apt-get install -y ntp
```

Set your network node to follow up your controller node

```
sed -i 's/server ntp.ubuntu.com/server 10.0.0.11/g' /etc/ntp.conf
```

Restart NTP service

```
service ntp restart
```

Install other services

```
apt-get install -y vlan bridge-utils
```

Edit /etc/sysctl.conf to contain the following

```
vi /etc/sysctl.conf
net.ipv4.ip_forward=1
```

```
net.ipv4.conf.all.rp_filter=0
net.ipv4.conf.default.rp_filter=0
```

Implement the changes

```
sysctl -p
```

Install the Networking components

```
apt-get install -y neutron-plugin-ml2 neutron-plugin-openvswitch-agent dnsmasq
neutron-l3-agent neutron-dhcp-agent
```

Update /etc/neutron/neutron.conf

```
vi /etc/neutron/neutron.conf

[DEFAULT]
auth_strategy = keystone
rpc_backend = neutron.openstack.common.rpc.impl_kombu
rabbit_host = 10.0.0.11
replace_core_plugin = neutron.plugins.ml2.plugin.Ml2Plugin with
core_plugin = ml2
service_plugins = router
allow_overlapping_ips = True

[keystone_authtoken]
auth_uri = http://10.0.0.11:5000
auth_host = 10.0.0.11
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = neutron
admin_password = service_pass
```

Edit the /etc/neutron/l3_agent.ini

```
vi /etc/neutron/l3_agent.ini

[DEFAULT]
interface_driver = neutron.agent.linux.interface.OVSInterfaceDriver
use_namespaces = True
```

Edit the /etc/neutron/dhcp_agent.ini

```
vi /etc/neutron/dhcp_agent.ini

[DEFAULT]
interface_driver = neutron.agent.linux.interface.OVSInterfaceDriver
dhcp_driver = neutron.agent.linux.dhcp.Dnsmasq
use_namespaces = True
```

Edit the /etc/neutron/metadata_agent.ini

```
vi /etc/neutron/metadata_agent.ini

[DEFAULT]
auth_url = http://10.0.0.11:5000/v2.0
auth_region = regionOne
```

```
admin_tenant_name = service
admin_user = neutron
admin_password = service_pass
nova_metadata_ip = 10.0.0.11
metadata_proxy_shared_secret = helloOpenStack
```

SSh to Controller VM and edit nova.conf

```
vi /etc/nova/nova.conf

[DEFAULT]
service_neutron_metadata_proxy = true
neutron_metadata_proxy_shared_secret = helloOpenStack

service nova-api restart
```

Now Switch back to Network VM and edit the following file /etc/neutron/plugins/ml2/ml2_conf.ini

```
vi /etc/neutron/plugins/ml2/ml2_conf.ini

[ml2]
type_drivers = gre
tenant_network_types = gre
mechanism_drivers = openvswitch

[ml2_type_gre]
tunnel_id_ranges = 1:1000

[ovs]
local_ip = 20.0.0.31
tunnel_type = gre
enable_tunneling = True

[securitygroup]
firewall_driver =
neutron.agent.linux.iptables_firewall.OVSHybridIptablesFirewallDriver
enable_security_group = True
```

Restart openVSwitch

```
service openvswitch-switch restart
```

Create the bridges:

```
#br-int will be used for VM integration
ovs-vsctl add-br br-int
```

```
#br-ex is used to make to VM accessible from the internet
ovs-vsctl add-br br-ex
```

Add the eth2 to the br-ex

#Internet connectivity will be lost after this step but this won't affect OpenStack's work

```
ovs-vsctl add-port br-ex eth2
```

Edit /etc/network/interfaces


```
vi /etc/network/interfaces

# The public network interface
auto eth2
iface eth2 inet manual
up ifconfig $IFACE 0.0.0.0 up
up ip link set $IFACE promisc on
down ip link set $IFACE promisc off
down ifconfig $IFACE down

auto br-ex
iface br-ex inet static
address 10.36.30.83
netmask 255.255.0.0
gateway 10.36.0.1
dns-nameservers 8.8.8.8
```

Restart network

```
ifdown eth2 && ifup eth2

ifdown br-ex && ifup br-ex
```

Restart all neutron services

```
service neutron-plugin-openvswitch-agent restart
service neutron-dhcp-agent restart
service neutron-l3-agent restart
service neutron-metadata-agent restart
service dnsmasq restart
```

Check status

```
service neutron-plugin-openvswitch-agent status
service neutron-dhcp-agent status
service neutron-l3-agent status
service neutron-metadata-agent status
service dnsmasq status
```

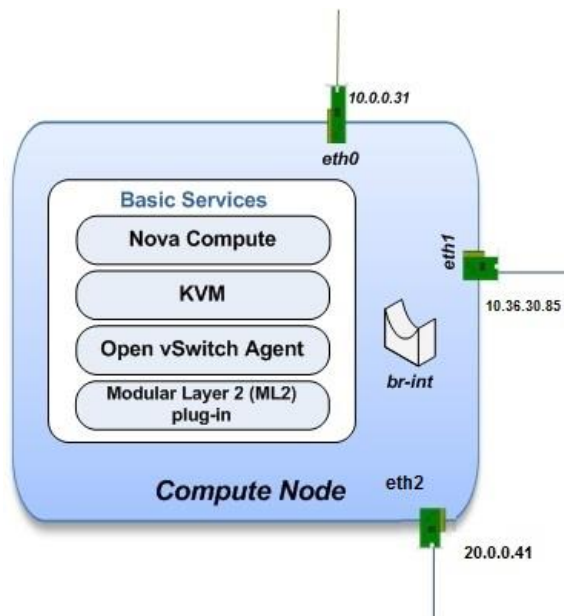
Create a simple credential file

```
vi creds
#Paste the following:
export OS_TENANT_NAME=admin
export OS_USERNAME=admin
export OS_PASSWORD=admin_pass
export OS_AUTH_URL="http://10.36.30.83:5000/v2.0/"
```

Check Neutron agents

```
source creds
neutron agent-list
```

Configure Compute Node



- **Update and Upgrade your System**

```
apt-get update -y && apt-get upgrade -y && apt-get dist-upgrade
```

- **Install ntp service**

```
apt-get install -y ntp
```

Set the compute node to follow up your controller node:

```
sed -i 's/server ntp.ubuntu.com/server 10.0.0.11/g' /etc/ntp.conf
```

- **Restart NTP service**

```
service ntp restart
```

- **Check that your hardware supports virtualization**

```
apt-get install -y cpu-checker  
kvm-ok
```

- **Install and configure kvm**

```
apt-get install -y kvm libvirt-bin pm-utils
```

- Install the Compute packages

```
apt-get install -y nova-compute-kvm python-guestfs
```

- Make the current kernel readable

```
dpkg-statoverride --update --add root root 0644 /boot/vmlinuz-$(uname -r)
```

- Enable this override for all future kernel updates, create the file `/etc/kernel/postinst.d/statoverride` containing

```
vi /etc/kernel/postinst.d/statoverride
#!/bin/sh
version="$1"
# passing the kernel version is required
[ -z "${version}" ] && exit 0
dpkg-statoverride --update --add root root 0644 /boot/vmlinuz-${version}
```

- Make the file executable

```
chmod +x /etc/kernel/postinst.d/statoverride
```

- Modify the `/etc/nova/nova.conf` like this

```
vi /etc/nova/nova.conf
[DEFAULT]
auth_strategy = keystone
rpc_backend = rabbit
rabbit_host = 10.0.0.11
my_ip = 10.0.0.31
vnc_enabled = True
vncserver_listen = 0.0.0.0
vncserver_proxyclient_address = 10.0.0.31
novncproxy_base_url = http://10.36.30.83:6080/vnc_auto.html
glance_host = 10.0.0.11
vif_plugging_is_fatal=false
vif_plugging_timeout=0

[database]
connection = mysql://nova:NOVA_DBPASS@10.0.0.11/nova

[keystone_authtoken]
auth_uri = http://10.0.0.11:5000
auth_host = 10.0.0.11
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = nova
admin_password = service_pass
```

- Delete `/var/lib/nova/nova.sqlite` file

```
rm /var/lib/nova/nova.sqlite
```

- Restart nova-compute services

```
service nova-compute restart
```

- Edit `/etc/sysctl.conf` to contain the following

```
vi /etc/sysctl.conf
net.ipv4.ip_forward=1
net.ipv4.conf.all.rp_filter=0
net.ipv4.conf.default.rp_filter=0
```

- **Implement the changes**

```
sysctl -p
```

- **Install the Networking components**

```
apt-get install -y neutron-common neutron-plugin-ml2 neutron-plugin-  
openvswitch-agent
```

- **Update /etc/neutron/neutron.conf**

```
vi /etc/neutron/neutron.conf

[DEFAULT]
auth_strategy = keystone
replace_core_plugin = neutron.plugins.ml2.plugin.Ml2Plugin with
core_plugin = ml2
service_plugins = router
allow_overlapping_ips = True
```

```
rpc_backend = neutron.openstack.common.rpc.impl_kombu
rabbit_host = 10.0.0.11
```

```
[keystone_authtoken]
auth_uri = http://10.0.0.11:5000
auth_host = 10.0.0.11
auth_port = 35357
auth_protocol = http
admin_tenant_name = service
admin_user = neutron
admin_password = service_pass
```

- **Configure the Modular Layer 2 (ML2) plug-in**

```
vi /etc/neutron/plugins/ml2/ml2_conf.ini

[ml2]
type_drivers = gre
tenant_network_types = gre
mechanism_drivers = openvswitch

[ml2_type_gre]
tunnel_id_ranges = 1:1000

[ovs]
local_ip = 20.0.0.41
tunnel_type = gre
enable_tunneling = True

[securitygroup]
firewall_driver =
neutron.agent.linux.iptables_firewall.OVSHybridIptablesFirewallDriver
enable_security_group = True
```

- Restart the OVS service

```
service openvswitch-switch restart
```

- Create the bridges

```
#br-int will be used for VM integration
ovs-vsctl add-br br-int
```

- Edit /etc/nova/nova.conf

```
vi /etc/nova/nova.conf

[DEFAULT]
network_api_class = nova.network.neutronv2.api.API
neutron_url = http://10.0.0.11:9696
neutron_auth_strategy = keystone
neutron_admin_tenant_name = service
neutron_admin_username = neutron
neutron_admin_password = service_pass
neutron_admin_auth_url = http://10.0.0.11:35357/v2.0
linuxnet_interface_driver =
nova.network.linux_net.LinuxOVSIInterfaceDriver
firewall_driver = nova.virt.firewall.NoopFirewallDriver
security_group_api = neutron

Edit /etc/nova/nova-compute.conf with the correct hypervisor type (set to
qemu if using virtualbox for example, kvm is default)

vi /etc/nova/nova-compute.conf

[DEFAULT]
compute_driver=libvirt.LibvirtDriver
[libvirt]
virt_type=qemu
```

- Restart nova-compute services

```
service nova-compute restart
```

- Restart the Open vSwitch (OVS) agent

```
service neutron-plugin-openvswitch-agent restart
```

- Check whether Nova is running

```
nova-manage service list
```

That's It ! Now your Openstack Cloud is UP and running.